

Anish Kshirsagar

DATA ENGINEER

Hyderabad, India (open to relocation — Bangalore) · anishkshirsagar02@gmail.com · +91-9908741411
linkedin.com/in/anish-kshirsagar02 · github.com/dante381

SUMMARY

Data Engineer with ~2 years building production-scale data infrastructure — high-volume ETL/ELT pipelines, lakehouse analytics, master data management, and data-reliability frameworks. Built an enterprise data-reliability platform from scratch that surfaces silent failures, sales drops, and duplicate records before they reach downstream consumers. Proven at MDM scale (80M+ records, 90%+ precision) across batch and real-time pipelines. Strong ownership mindset with a track record of improving data accuracy, pipeline performance (30–40% latency reduction), and data quality in production.

CORE SKILLS

Languages: Python, SQL, Java, C++, Node.js

Big Data: Apache Spark / PySpark, distributed batch processing, ETL/ELT

Cloud & Platforms: Azure Data Factory, Synapse, Microsoft Fabric, Databricks / Delta Lake, lakehouse architecture

Data Engineering: data modeling, partitioning, performance tuning, batch & real-time pipeline design

Data Reliability: schema validation, data-quality frameworks, anomaly detection, monitoring, CI/CD basics

ML Data Systems: embeddings, similarity search, probabilistic record linkage, clustering

EXPERIENCE

Data Engineer 1 — Data & Analytics, Northern Tool + Equipment

May 2024 – Present

- Owned end-to-end Azure cloud data pipelines processing customer data from 5+ heterogeneous sources (CRM, ERP, third-party feeds) into a centralized lakehouse — designing both batch and real-time architecture to meet SLA and reliability requirements.
- Built an enterprise data-reliability platform with automated schema enforcement, null validation, and anomaly detection — reducing downstream data issues by **35%** and surfacing hidden failures (sales drops, duplicate indexes, silent pipeline faults) before business impact.
- Designed a Master Customer Dataset using probabilistic record linkage, deduplicating **80M+ records to 40M unique customers** at **90%+ precision**.
- Built scalable batch embedding pipelines on Spark (PySpark), vectorizing millions of text and categorical attributes — reduced false matches by **25%** via cosine-similarity evaluation workflows.
- Optimized PySpark configurations, SQL transformations, and partitioning strategies, cutting pipeline runtime by **30–40%** across production workloads.
- Developed data modeling and partitioning aligned with Delta Lake lakehouse architecture, enabling efficient query patterns and incremental processing.

Software Developer Intern, Northern Tool + Equipment

Sept 2023 – Apr 2024

- Built a Flutter-based real-time telemetry application ingesting BLE data from Northstar generators.
- Reduced data ingestion and connection latency by **40%** through optimized batching and background processing; contributed to 95%+ user satisfaction.

Teaching Assistant, ProCareer Academy

Aug 2023 – Sept 2023

- Mentored 100+ students in Python and backend fundamentals via hands-on labs and debugging sessions.

PROJECTS

PC Monitoring Platform — Distributed PC monitoring system (Unreal Engine + backend services); reduced service-desk workload by 25%. Multi-language (C++, Python, Java) modular ingestion and observability.

Virtual Assistant Chatbot — Full-stack application integrating OpenAI APIs; scaled experimentation to 30% more active users.

Smart Home Security System — IoT + AI security platform with real-time ingestion and alerting; reduced false alarms by 35%.

EDUCATION

B.Tech in Computer Science — Keshav Memorial Institute of Technology (KMIT) · 2020 – 2024