

Anish Kshirsagar

+91-9908741411 | anishkshirsagar02@gmail.com | linkedin.com/in/anish-kshirsagar02 | github.com/dante381 |
Hyderabad, India

PROFESSIONAL SUMMARY

Data Engineer with end-to-end ownership of production ETL pipelines across heterogeneous data sources (CRM, ERP, third-party feeds), delivering analytics-ready datasets to cross-functional Analytics, ML, and reporting teams. Implemented data quality automation reducing downstream issues by 35%; probabilistic record linkage achieved 90%+ precision at 80M+ record scale. Spark and distributed batch processing are primary production tools. Azure Synapse and Microsoft Fabric experience parallels BigQuery ecosystem paradigms. Embedding pipeline work at an AI company context is directly relevant to Cohere's data foundations mission.

CORE COMPETENCIES

ETL/ELT

Apache Spark (PySpark)

Data Modeling

Cross-functional Delivery

Python

SQL

Data Quality Frameworks

Lakehouse Architecture

Azure Synapse Analytics

Microsoft Fabric

Record Linkage

Performance Tuning

Schema Validation

Embedding Pipelines

WORK EXPERIENCE

Northern Tool + Equipment May 2024 - Present

Data Engineer 1 - Data & Analytics

- Owned end-to-end Spark-based ETL pipelines processing customer data from 5+ heterogeneous sources (CRM, ERP, third-party feeds) into a centralized lakehouse - full pipeline ownership from ingestion through transformation to analytics-ready delivery.
- Implemented automated data quality frameworks (schema enforcement, null validation, anomaly detection), reducing downstream data issues by 35% and establishing reliable data foundations for analytics and ML stakeholders.
- Partnered with analytics, ML, and reporting teams across the organization to define data requirements and deliver analytics-ready datasets supporting customer intelligence, reporting, and experimentation.
- Designed and implemented a Master Customer Dataset using probabilistic record linkage, improving identity resolution from fragmented records to 90%+ precision across 80M+ records - demonstrates operating at the edge of standard ETL.
- Built scalable batch embedding pipelines on Spark, transforming millions of text and categorical attributes into vectorized features - ML data background directly relevant to Cohere's LLM-adjacent data needs.
- Optimized SQL transformations, Spark job configurations, and partitioning strategies, cutting pipeline runtime by 30-40%.
- Developed embedding evaluation workflows using cosine similarity and distance metrics, reducing false matches by 25% and improving downstream data quality.

Northern Tool + Equipment Sept 2023 - Apr 2024

Software Developer Intern

- Built a Flutter-based real-time telemetry application ingesting BLE data from Northstar generators.
- Reduced data ingestion and connection latency by 40% through optimized batching and background processing.
- Improved system reliability and user experience, contributing to 95%+ user satisfaction.

ProCareer Academy Aug 2023 - Sept 2023

Teaching Assistant

- Mentored 100+ students in Python and backend fundamentals via hands-on labs and debugging sessions.
- Designed practical assignments that improved problem-solving depth and engagement by 40%.

PROJECTS

PC Monitoring Platform [GitHub](#)

Built a distributed PC monitoring system using Unreal Engine and backend services, reducing service-desk workload by 25%. Integrated multi-language components (C++, Python, Java) for modular data ingestion and observability.

Virtual Assistant Chatbot [GitHub](#)

Developed a full-stack application integrating OpenAI APIs, enabling scalable AI experimentation for 30% more active users.

Smart Home Security System [GitHub](#)

Built an IoT + AI security platform with real-time data ingestion and alerting, reducing false alarms by 35%.

EDUCATION

B.Tech in Computer Science - [Keshav Memorial Institute of Technology](#) Sept 2020 - May 2024

SKILLS

Data Engineering: ETL/ELT, Data Modeling, Partitioning, Performance Tuning

Big Data: Apache Spark (PySpark), Distributed Batch Processing **Programming:** Python, SQL, Java, C++, NodeJS

Data Platforms: Lakehouse Architecture, Microsoft Fabric, Azure Synapse Analytics

ML Data Systems: Embeddings, Similarity Search, Clustering, Record Linkage

Data Reliability: Schema Validation, Data Quality Checks, Monitoring, Git, CI/CD Basics